

***F1000Research* Evaluating Metagenome Assembly on a Complex Community**

Sherine Awad¹, Luiz Irber², and C. Titus Brown³

^{1,2,3}Department of Population Health and Reproduction, University of California, Davis, California

Abstract

Metagenome assembly is a challenging problem due to the biodiversity of the microorganisms. Most assemblers are designed for whole genome assembly and not capable of dealing with metagenomic samples. However, in order to decide which assembler works best for metagenome, we need to evaluate metagenome assembly generated by each assembler.

In this paper, we used three assemblers ; IDBA-UD, SPAdes, and Megahit to assemble metagenome mock community data and evaluate the assembly process in terms of resources utilization, assembly quality, genome fraction covered, duplication ratio, misassemblies and partial alignments.

The results show only small differences in content recovery between assemblers. However, Megahit is much faster and produces shorter contig lengths than IDBA-UD and SPAdes.

Introduction

Metagenomics refers to sequencing of DNA from a mixture of organisms, often from an environmental or uncultured sample. Unlike whole genome sequencing, metagenomics targets a mixture of genomes, which introduces metagenome-specific challenges in analysis. Most approaches to analyzing metagenomic data rely on mapping or comparing sequencing reads to reference sequence collections. However, reference databases contain only a small subset of microbial diversity (cite: geba), and the much of the remaining diversity is evolutionarily distant and search techniques may not access it.

As sequencing capacity increases and sequence data is generated from many more environmental samples, metagenomics is increasingly using de novo assembly techniques to generate new reference genomes and metagenomes. There are a number of metagenome assemblers that are widely used. However, evaluating the results of these assemblers is challenging due to the general lack of good quality reference metagenomes. Below, we evaluate three commonly assemblers - SPAdes, IDBA, and MEGAHIT - on a mock community containing 64 species of microbes with known genomes.

Moya et al. in [?] evaluated metagenome assembly using two simulated 454 viral metagenome and six assemblers. The assemblies were evaluated based on several metrics including N50, percentages of reads assembled, accuracy when compared to the reference genome. In addition to, chimeras per contigs and the effect of assembly on taxonomic and functional annotations.

Mavromatis et al. in [?] provided a benchmark study to evaluate the fidelity of metagenome process methods. The study used simulated metagenomic data sets constructed at different complexity levels. The datasets were assembled using Phrap v3.57, Arachne v.2 [?] and JAZZ. [?] This study evaluates assembly, gene prediction, and binning methods. However, the study did not evaluate the assembly quality against a reference genome.

Rangwala et al. in [?] presented an evaluation study of metagenome assembly. The study used a de Bruijn graph based assembler ABYSS [?] to assemble simulated metagenome reads of 36 bp. The data set is classified at different complexity levels. The study compares the quality of the assembly of the data sets in terms of quality measures of contigs length, assembly accuracy. The study also took into consideration the effect of kmer size and the degree of chimericity. However, the study evaluated the assembly based on one assembler, and did not evaluate assembly against several assemblers. Also, both previous studies used simulated data, which may lack confounders of assembly such as sequencing artifacts and GC bias.

Shakya et al. (2013) constructed a complex synthetic community of organisms by mixing DNA isolated from individual cultures of 64 bacteria and archaea. In addition to performing 16s amplicon analysis and doing 454 sequencing, the authors shotgun sequenced the mixture with Illumina (@cite). While the authors concluded that this metagenomic sequencing generally out-

performed amplicon sequencing, they conducted a mapping based analysis rather than an assembly based analysis.

More recently, several benchmark studies systematically evaluated metagenome assembly of short reads. (CAMI and ...)

In this paper, we evaluate metagenome assembly on the Illumina data set from Shakya et al. (2013) using three assemblers; IDBA-UD [?], SPAdes [?], and MEGAHIT [?]. These three assemblers were chosen because they are actively used and highly cited.

SPAdes [?] is an assembler for both single-cell and standard (multicell) assembly. (More description here @CTB.)

IDBA-UD [?] is a de Bruijn graph approach for assembling reads from single cell sequencing or metagenomic sequencing technologies with uneven sequencing depths. IDBA-UD uses multiple depth-relative thresholds to remove erroneous k-mers in both low-depth and high-depth regions. It also uses paired-end information to solve the branch problem of low-depth short repeat regions. It also applies an error correction step to correct reads of high-depth regions that can be aligned to high confident contigs.

MEGAHIT [?] is a newer approach that constructs a succinct de Bruijn graph using multiple k-mer sizes, and uses a novel “mercy k-mer” approach that preserves low-abundance regions of reads. It also can use GPUs to accelerate the graph construction.

Below, we evaluate the performance of these three assemblers using the synthetic (“mock”) community data from the Shakya et al. study. The performance of each assembler is compared in terms of resource utilization, covered genome fraction, duplication ratio, gene recovery, contig misassembly, and contig length.

In contrast to some other evaluations (assemblathon 2), we work with a synthetic community with a known answer, do not tune the default parameters, and investigate memory and runtime as a key part of the process. This mimics default user experience. (Expand.)

In this report, we update the list of reference genomes for Shakya et al. to include plasmids and compare IDBA, SPAdes, and MEGAHIT performance on short-read metagenome assembly. We also evaluate inter-strain chimerism in the assemblies, explore the failure to recover a specific *Shewanella baltica* strain variant in the assembly, and detect several previously unreported genomes in the Shakya et al. data set. Our report provides strong guidance on choice of assemblers and significantly extends previous analyses of this low-complexity metagenome benchmarking data set.

Datasets

We used a diverse mock community data set constructed by pooling DNA from 64 species of bacteria and archaea and sequencing them with Illumina HiSeq. The raw data set consisted of 109,629,496 reads from Illumina HiSeq 101 bp paired-end sequencing (2x101) with an




Figure 1. Cumulative coverage profile for the reference metagenome, based on read mapping.

untrimmed total length of 11.07 Gbp and an estimated fragment size of 380 bp [?].

The original reads are available through the NCBI Sequence Read Archive at Accession SRX200676. We received the 64 reference genomes from the original authors. They consist of 205.6 Mbp of assembled genomes in 64 contigs, and are available for download at <https://dx.doi.org/10.6084/m9.figshare.1506873.v2>. We updated the data sets from NCBI etc. etc. The following genomes were updated. Updated data is available for download here (OSF).

Methods

The analysis code and run scripts for this paper are available at: <https://github.com/dib-lab/2015-metagenome-assembly/>. The scripts and overall pipeline were examined by the first and senior authors for correctness. In addition, the bespoke reference-based analysis scripts were tested by running them on a single-colony *E. coli* MG1655 data set with a high quality reference genome [?].

Quality Filtering

We removed adapters with Trimmomatic v0.30 in paired-end mode with the Truseq adapters [?], using light quality score trimming as recommended in MacManes, 2014 [?].

Reference Coverage Profile

To evaluate how much of the reference metagenome was contained in the read data, we used `bwa aln` to map reads to the reference genome. We then calculated how many reference bases were covered by how many mapped reads (custom script `coverage-profile.py`).

Assemblers

We assembled the quality-filtered reads using three different assemblers: IDBA-UD [?], MetaSPAdes [?], and MEGAHIT [?]. For IDBA-UD v1.1.1 [?], we used

`-pre_correction` to perform pre-correction before assembly and `-r` for the `pe` files.

For MetaSPAdes v3.9.0 [?], we used `-meta` `-pe1-12` `-pe1-s` where `-meta` is used for metagenomic data sets, `-pe1-12` specifies the interlaced reads for the first paired-end library, and `-pe1-s` provides the orphan reads remaining from quality trimming.

For MEGAHIT v1.1.1-2-g02102e1 [?], we used `-l 101` `-m 3e9` `-cpu-only` where `-l` is for maximum read length, `-m` is for max memory in bytes to be used in constructing the graph, and `-cpu-only` to use only the CPU and no GPUs. We also used `-presets meta-large` for large and complex metagenomes, and `-12` and `-r` to specify the interleaved-paired-end and single-end files respectively. MEGAHIT allows the specification of a memory limit and we used `-M 1e+10` for 10 GB.

All three assemblies were executed on the same high-memory buy-in node on the Michigan State University High Performance Compute Cluster, and we recorded RAM and CPU time of each assembly job using the `qstat` utility at the end of each run.

Unless otherwise mentioned, we eliminated all contigs less than 500 bp from each assembly prior to further analysis.

Mapping

We aligned all quality-filtered reads to the reference metagenome with `bwa aln` (v0.7.7.r441) [?]. We aligned paired-end and orphaned reads separately. We then used `samtools` (v0.1.19) [?] to convert SAM files to BAM files for both paired-end and orphaned reads. To count the unaligned reads, we included only those records with the “4” flag in the SAM files [?].

To extract the reads that contribute to unaligned contigs, we mapped the quality filtered reads to the unaligned contigs using `bwa aln` (v0.7.7.r441) [?]. Then we used `samtools` to retrieve the reads that mapped to the unaligned contigs.

k-mer Presence

In order to examine k-mer presence for a k-mer size of 20, we built a k-mer counting table from the given quality filtered reads using `load-into-counting.py` from `khmer` [?]. Then we calculate abundance distribution of the k-mers in the quality filtered reads using the pre-made k-mer counting table using `abundance-dist.py`. We followed the same approach to examine k-mer presence in assemblies.

Assembly analysis using Nucmer

We used the NUCmer tool from MUMmer3.23 [?] to align assemblies to the reference genome with options `-coords` `-p`. Then we parsed the generated “coords” file using a custom script `analyze_assembly.py`, and calculated several analysis metrics across all three assemblies at two alignment identities, 95% and 99%.

Reference-based analysis of the assemblies

We analyzed metrics for three different sets of contigs, based on the NUCmer alignments. We used the unfiltered NUCmer alignments for the analyses termed “ambiguous.” We also subjected the alignments to two different filtering criteria, “best-hit” and “no-misassemblies.” In the best-hit approach, among all alignments of a contig, we took into consideration the longest alignment with an identity above a specified identity threshold (either 95% or 99%). In the no-misassemblies approach, we only counted contigs that have precisely one alignment within the reference.

In all approaches, we flag a base in the reference genome as “covered” if it is contained in a kept alignment. We define the duplication ratio as the percentages of bases in the reference covered by two or more kept alignments. We define misassemblies as those contigs that are divided into different parts when mapped to the reference. The number of misassembled contigs is equal to the number of aligned contigs (both totally and partially) in the ambiguous approach, minus the number of aligned contigs in the no-misassemblies approach.

All approaches have a non-zero duplication ratio within the reference because we do not explicitly discard contigs that map to the same location in the reference.

Analysis of chimeric misassemblies

We analyzed each assembly for chimeric misassemblies by counting the number of contigs that contained matches to two distinct reference genomes. In order to remove secondary alignments from consideration, we first filtered matches extracted the longest non-overlapping NUCmer alignments for each contig at a minimum alignment identity of 99%. We then used the script `analyze_chimeric2.py` to find contigs that matched sequences two or more distinct reference species. As a negative control on our analysis, we verified that this approach yielded no positive results when applied to the alignments of the reference metagenome against itself.

Results

The raw data is high quality.

The reads contains 11,072,579,096 bp (11.07 Gbp) in 109,629,496 reads with 101.0 average length (2x101bp Illumina HiSeq).

Trimming removed 686,735 reads (0.63%). After trimming, we retained 108,422,358 paired reads containing 10.94 Gbp with an average length of 100.9 bases. A total of 46.56 Mbp remained in 520,403 orphan reads with an average length of 89.5 bases. In total, the quality trimmed data contained 10.98 Gbp in 108,942,761 reads. This quality trimmed (“QC”) data set was used as the basis for all further analyses.

Table 1. Jaccard containment of the reference in the reads

k-mer size	% reference in reads
21	96.8%
31	95.9%
41	94.9%
51	94.1%

The reference metagenome is not completely present in the reads.

We next evaluated the fraction of the reference genome covered by at least one read (see Methods for details). Quality filtered reads cover 203,058,414 (98.76%) bases of the reference metagenome (205,603,715 bp total size). Figure 1 shows the cumulative coverage profile of the reference metagenome, and the percentage of bases with that coverage. Most of the reference metagenome was covered at least minimally; only 3.33% of the reference metagenome had mapping coverage <5, and 1.24% of the bases in the reference were not covered by any reads in the QC data set.

In order to evaluate reconstructability with De Bruijn graph assemblers, we next examined k-mer containment of the reference in the reads for k of 21, 31, 41, and 51 (Table 1). The k-mer overlap decreases from 96.8% to 94.1% as the k-mer size increases. This could be caused by low coverage of some portions of the reference and/or variation between the reads and the reference.

Some individual reference genomes are poorly represented in the reads.

Table 2. Top uncovered genomes

Genome	Read coverage	21-mer presence
<i>B. xenovorans</i>	99.2%	91.6%
<i>D. vulgaris</i> DP4	91.98%	79.1%
<i>T. thermophilus</i> HB27	91.06%	79.7%
<i>E. faecalis</i> V583	78.14%	68.0%
<i>F. nucleatum</i>	47.6%	17.8%
	89.22%	XX%
	87.69%	XX%
	85.16%	XX%
	80.00%	XX%
	79.07%	XX%
<i>E. faecalis</i>	74.62%	XX%
<i>D. vulgaris</i>	62.37%	XX%
<i>E. faecalis</i>	61.86%	XX%
<i>D. radiodurans</i>	55.33%	XX%
<i>F. nucleatum</i>	47.58%	XX%

To see if specific reference genomes exhibited low coverage, we analyzed read mapping coverage and 21-mer containment for individual genomes. Of the 64 reference genomes used in the metagenome, 59 had a per-base mapping coverage above 95% and a 21-mer containment

in the QC reads above 95%. The remaining five varied significantly in both metrics (Table 3), with *E. nucleatum* the lowest – only 47.30% of the bases in the reference genome are covered by one or more mapped reads, and only 17.8% of the 21-mers in the *E. nucleatum* reference genome are present in the reads at any abundance.

We next did a 51-mer containment analysis of each reference genome in the reads. 99% or more of the constituent 51-mers for 51 of the 64 reference genomes were present in the reads, suggesting that each of the 51 genome was entirely present at some minimal coverage.

We excluded the remaining 13 genomes from any comparative analysis of assembly quality, because interpreting coverage and misassembly analysis for these genomes would be impossible. (@CTB list or table?)

MEGAHIT is the fastest and lowest-memory assembler evaluated

Table 3. Running Time and Memory Utilization

Assembler	CPU time	Wall time	RAM
MEGAHIT	52hr 25m	4 hr 9m	11.4 GB
IDBA-UD	17h		149.1 GB
SPAdes	94hr 43m	94hr 44m	100.7 GB

We ran three commonly used metagenome assemblers on the QC data set: IDBA-UD, SPAdes, and MEGAHIT. We recorded the time and memory usage of each (Table 3). In computational requirements, MEGAHIT outperformed both SPAdes and IDBA-UD considerably, producing an assembly in four hours – approximately 4 times faster than IDBA and 8 times faster than SPAdes. MEGAHIT used only 11.4 GB of RAM – 1/13th to 1/9th the memory used by IDBA and SPAdes, respectively.

The assemblies contain most of the raw data

Table 4. Read and k-mer exclusion from assemblies

Assembly	Unmapped Reads	51-mers omitted
IDBA	3,328,674 (3.05%)	3.4%
SPAdes	3,879,573 (3.56%)	4.2%
MEGAHIT	5,848,494 (5.37%)	3.7%

We assessed read inclusion in assemblies by mapping the QC reads to the length-filtered assemblies and counting the remaining unmapped reads. Depending on the assembly, between 3.3 million and 5.9 million reads (3.0-5.4%) did not map to the assemblies (Table 4). Here, the MEGAHIT assembly was distinguished by representing 2 million fewer reads than the IDBA and SPAdes assemblies. K-mer inclusion, however, was more closely matched across the assemblies, with all three assemblies contain-

ing between 95.8% and 96.6% of the 51-mers in the k-mer trimmed reads.

Much of the reference is covered by the assemblies.

Table 5. Contig coverage of reference with “loose” alignment conditions.

Assembly	% covered	% Duplication
MEGAHIT	96.2%	0.72%
SPAdes	95.8%	0.99%
IDBA	95.6%	0.88%

(Update this section to include k-mer overlap - see comments.)

We next evaluated the extent to which the assembled contigs recovered the “known/true” metagenome sequence by aligning each assembly to the adjusted reference (Table 5). Each of the three assemblers generates contigs that cover more than 95.6% of the reference metagenome at high identity (99%) with little duplication (0.72-0.99%). At 99% identity with the loose mapping approach, approximately 1.8% of the reference is missed by all three assemblers, while 0.9% is uniquely covered by MEGAHIT, 0.6% is uniquely covered by SPAdes, and 0.4% is uniquely covered by IDBA.

The generated contigs are broadly accurate.

Table 6. Contig accuracy measured by reference coverage with strict alignment.

Assembly	% covered
MEGAHIT	93.8%
IDBA	89.5%
SPAdes	87.3%

When counting only the best (longest) alignment per contig at a 99% identity threshold, each of the three assemblies recovers more than 87.3% of the reference, with MEGAHIT recovering the most – 93.8% of the reference (Table 6).

(CTB: add discussion of accuracy/mismatches/indels here.)

Assembly statistics for individual genomes.

The NGA50 for a particular genome in a metagenome is the minimum contig size at which 50% of that genome is covered by aligned contigs. We computed the NGA50 for each individual genome and assembly (see Figure ??). The NGA50 statistics for individual genomes vary widely, but there are consistent assembler-specific trends: MEGAHIT yields the lowest NGA50 for 42 of the 46 genomes, while SPAdes yields the highest NGA50 for 38 of the 46 genomes.

Longer contigs are less likely to be chimeric.

Table 7. Chimeric contigs by contig length.

Assembly	> 50kb	> 5kb	> 500 bp
MEGAHIT	0%	3.9%	15.5%
IDBA	0%	6.3%	22.2%
SPAdes	0%	8.8%	18.4%

Chimerism is the formation of contigs that include sequence from multiple genomes. We evaluated the rate of chimerism in contigs at three different contig length cutoffs: 500bp, 5kb, and 50kb (Table ??). We found that the percentage of contigs that match to the genomes of two or more different species drop as the minimum contig size increases, to the point where contigs longer than 50kb had no chimeric contigs in any assembly. Upon further analysis, 99% of the chimeric contigs greater than 5kb were formed between genomes from pairs of closely related species. Similarly, 80% of the chimeric contigs at the lowest cutoff (500 bp) were between pairs or groups of closely related species.

The unmapped reads contain strain variants of reference genomes.

Approximately 7.6 million reads (7.00%) from the QC data set did not map anywhere in the reference provided by the authors of @cite. We extracted and assembled these reads in isolation using MEGAHIT, yielding XXX bases YYY. We then did a k-mer analysis of this assembly against all of the RefSeq genomes at k=51 using sourmash @cite, and estimated the fraction of the contigs that belonged to different species (Table ??). We estimate that 50.3% of the contigs have > 1% k-mer content that positively matches to a genome not present in the reference metagenome (although note *Shewanella*).

To verify these assignments, we aligned the assembly of unmapped reads to these RefSeq genomes in aggregate under the “loose” alignment criteria and found XXX.

We note that all but one of the matches in Table ??, *P. ruminis*, are from genomes of strain variants of putative members of the reference metagenome. Moreover, 4 of the 5 top uncovered reference genomes in the reads (Table 2) are from strain variants we find in the assembly of the unmapped reads – only *B. xenovorans* is omitted.

The majority of the matches are to strain variants of species present in the mock community design: e.g. the *Nostoc* sp. present in the reads is 87.8% similar by Jaccard at k=31 to the *Nostoc* sp. in the mock, the *Ruegeria pomeroyi* strains are 87.6% similar at k=31, and the *Herpetosiphon aurantiacus* strains are 92.4% similar at k=31. (Thermus as well.) There are many *Fusobacterium* matches as well.

Interestingly, of these, a few of the mock community members are incomplete in the reads (*Fusobacterium*, *Ruegeria*, *Thermus*, *Enterococcus*) while others are entirely present (*Nostoc*, *Herpetosiphon*). *Nostoc* and *Herp*: plasmids that aren’t present in the original accession.

Table 8. RefSeq strains detected in MEGAHIT assembly of unmapped reads

Fraction	RefSeq genome
9.6%	NC_013968.1 <i>H. volcanii</i> DS2 plasmid pHV1
8.0%	NZ_LN831027.1 <i>F. nucleatum</i> polymorphum
7.4%	NC_003272.1 <i>Nostoc</i> sp. PCC 7120
5.3%	NC_003911.12 <i>R. pomeroyi</i> DSS-3
3.2%	NZ_CH959311.1 <i>Sulfitobacter</i> sp. EE-36
3.0%	NC_006461.1 <i>T. thermophilus</i> HB8
2.4%	NZ_JNKC01000001.1 <i>P. ruminis</i> DSM 24773
2.3%	NZ_CH959317.1 <i>Sulfitobacter</i> sp. NAS-14.1
1.9%	NC_011663.1 <i>S. baltica</i> OS223
1.8%	NZ_KI965381.1 <i>F. nucleatum</i> 13_3C
1.6%	NC_002937.3 <i>D. vulgaris</i> str. Hildenborough
1.4%	NC_001263.1 <i>D. radiodurans</i> R1
1.3%	NZ_AFHH01000001.1 <i>E. faecalis</i> OG1X
1.1%	NZ_AXNV01000001.1 <i>F. nucleatum</i> CTI-6
50.3%	(Sum fraction of contigs)

Should we add them?? Yeah probably.

The presence of about 700 KB of *Proteiniclasticum ruminis* is a puzzle, since there is no closely related species in the mock community design. Note that very little of the assembly aligns to the known *P. ruminis* genome via nucmer at 99%, suggesting that this is a strain variant (most alignments are at 94% or higher, representing 2.73 Mbp).

Shewanella baltica OS223 presents a puzzling situation. While *Shewanella baltica* OS223 was removed from the reference metagenome for having less than 99% of 51-mers present in the reads, the genome is well represented in the read data set - 99.95% of the genome is covered by at least one read. We would therefore expect to recover a considerable fraction of the OS223 genome in the reads that didn’t map to the reference metagenome; however, we find less than XXX bases in the MEGAHIT assembly that align to OS223 at 99% or higher identity. Why is OS223 is present but doesn’t assemble? This appears to be due to the presence of a closely related strain, *Shewanella baltica* OS185. Despite a fairly low Jaccard distance of 37.5% at k=31, approximately 2/3 of the QC reads that map to OS223 also map to OS185 with our choice of parameters to `bwa aln`, while less than 5% of the OS223 genome aligns to OS185 with nucmer at 99% or higher. We infer that the assemblers are “choosing” to assemble OS185 based on its higher abundance in the reads, and discarding OS223 reads as erroneous variants of OS185. This highlights another challenge presented by the presence of multiple strains in a sample: strains may assemble quite poorly (also observed in CAMI).

`./summarize-coords.py os223-vs-genomes/63.fa`
`NC_011663.1, 2501017NC_011664.1, 53226NC_011668.1, 45752NC_011669.1`
so about half of OS223 is covered in the regular assembly
what happens to OS185? about 80
maybe compute coverage stats for all genomes wrt assembly.

note plasmids covered!?

@CTB how many reads in unmapped map to shewanella?

@CTB how many reads map to both OS223 and OS185?
apprx 2/3.

@CTB check to see what's going on with nucmer align of
OS223 to OS185. (very little)

@CTB split to discussion?

@CTB abundance of one vs other in original?

Discussion

Assembly recovers basic content sensitively and accurately.

All three assemblers performed well in assembling contigs from the content that was fully present in reads and k-mers. After length filtering, all three assemblies contained more than 95% of the reference (Table 5); even with removal of secondary alignments, more than 87% was recovered by each assembler (Table 6). About half the constituent genomes had an NGA50 of 50kb or higher (Table ??), which, while low for current Illumina single-genome sequencing @cite, is sufficient to recover operon-level relationships for many genes.

Chimeric contigs form between closely related species.

A primary interest of this study is to examine chimerism, in which a single contig is formed from multiple genomes. Naively we expected chimerism to behave like misassembly, in which longer contigs would be more likely to contain large-scale errors. Surprisingly, we find less chimerism the longer the contigs were; the shorter contigs contained a high rate of chimerae, while contigs longer than 50kb had none (Table ??).

Analysis of the genomic pattern of chimeric contigs provides a possible explanation: the majority of chimeric contigs at all sizes are formed between closely related species or strain variants, suggesting that chimeric contigs are generated primarily where multiple species overlap in the assembly graph. However, the same complexity in the local assembly graph that leads to chimera may also limit the extension of contigs, resulting in shorter chimeric contigs. Regardless, we find that chimeric contigs form between closely related species with a high frequency.

MEGAHIT performs best by several metrics.

MEGAHIT is clearly the most efficient computationally, outperforming both SPAdes and IDBA by 5-10x in memory and 17-42x in time (Table 3). While the MEGAHIT assembly had 2m fewer reads mapping to it than the other assemblies (Table 1), MEGAHIT covered more of the reference genome with both loose and strict alignments (Table 5 and Table 6), with little duplication.

While MEGAHIT yields the lowest NGA50 of the three assemblers on this data set, MEGAHIT also had the lowest rate of chimerism. As discussed above, this may be an un-

avoidable tradeoff; either way, for complex populations, we believe conservative assembly is a feature rather than a "bug".

Between the assemblers, the assembly content differs by only a small amount when loose alignments are allowed: all three assemblers miss more content (approximately 1.8% of the reference) than they generate uniquely (0.9% or less). In addition to preferring no one assembler over any other, this suggests that combining assemblies may have little value in terms of recovering additional metagenome content.

The missing reference may be present in strain variants of the intended species.

Several individual genomes are missing in measurable portion from the QC reads (Table 2), and many QC reads (7% of 108m) did not map to the provided reference genome. These may be related issues: upon analysis of the unmapped reads against RefSeq, we find that about half of the contigs assembled from the unmapped reads can be assigned to strain variants of the species in the mock community (Table ??). This suggests that the mock community may have inadvertently included strain variants of the intended microbes. We tentatively confirm this by doing a k-mer analysis of the 454 data from @cite against the assembled contigs and showing that XXX.

Without returning to the original DNA samples, it is impossible to conclusively confirm that unintended strains were used in the construction of the mock community. In particular, our analysis is dependent on the genomes in RefSeq: the genomes we detect in the contigs are clearly more closely related to the contigs than the species in the reference metagenome, based on mapping, k-mer analysis, and contig alignment; but RefSeq is unlikely to contain the exact genomes of the included strain variants.

Conclusions

Overall, assembly of this moderately complex mock community works well, with good recovery of known genomic sequence with few inaccuracies or misassemblies. All three assemblers that we evaluated recover similar amounts of sequence with similar error rates, although MEGAHIT is computationally most efficient. MEGAHIT also clearly produces consistently shorter NGA50s and the fewest chimeric contigs; we believe these two facts reflect an important conservatism in MEGAHIT's contig assembly algorithm. We can therefore unambiguously recommend MEGAHIT over SPAdes and IDBA, based on both sequence recovery and the lower rates of chimerism, as well as MEGAHIT's dramatically better computational performance.

A more general conclusion from our analysis is that shorter contigs are considerably more likely to be chimeric misassemblies than long, in even moderately complex communities; here, this is based on positive identification of the chimerae from good quality references. We infer that strain variation may well be the cause of many

short contigs. We also note that, in the absence of high-quality mock communities or isolate reference genomes, chimerism between strain variants is extremely hard to evaluate, and this trend may have gone undetected in studies of real microbial communities.

An additional concern is that metagenome assemblies are often performed after pooling data sets to increase coverage; this pooled data is however more likely to contain multiple strains, which would then in turn adversely affect assembly of strains. This may not be resolvable within the current paradigm of assembly, which focuses on outputting linear assemblies that cannot represent strain variation.

We next plan to examine this trend in other mock communities, and with longer reads. (ref Disco)

Author contributions

In order to give appropriate credit to each author of an article, the individual contributions of each author to the manuscript should be detailed in this section. We recommend using author initials and then stating briefly how they contributed.

Competing interests

No competing interest to our knowledge.

Grant information

This work is funded by Moore and NIH.

Acknowledgments

Michael R. Crusoe

References

- [] Jorge F Vázquez-Castellanos, Rodrigo García-López, Vicente Pérez-Brocal, Miguel Pignatelli, and Andrés Moya. Comparison of different assembly and annotation tools on analysis of simulated viral metagenomic communities in the gut. *BMC genomics*, 15(1):1, 2014.
- [] Konstantinos Mavromatis, Natalia Ivanova, Kerrie Barry, Harris Shapiro, Eugene Goltsman, Alice C McHardy, Isidore Rigoutsos, Asaf Salamov, Frank Korzeniewski, Miriam Land, et al. Use of simulated data sets to evaluate the fidelity of metagenomic processing methods. *Nature methods*, 4(6):495–500, 2007.
- [] David B Jaffe, Jonathan Butler, Sante Gnerre, Evan Mauceli, Kerstin Lindblad-Toh, Jill P Mesirov, Michael C Zody, and Eric S Lander. Whole-genome sequence assembly for mammalian genomes: Arachne 2. *Genome research*, 13(1):91–96, 2003.
- [] Samuel Aparicio, Jarrod Chapman, Elia Stupka, Nik Putnam, Jer-ming Chia, Paramvir Dehal, Alan Christoffels, Sam Rash, Shawn Hoon, Arian Smit, et al. Whole-genome shotgun assembly and analysis of the genome of fugu rubripes. *Science*, 297(5585):1301–1310, 2002.
- [] Anveshi Charuvaka and Huzefa Rangwala. Evaluation of short read metagenomic assembly. *BMC genomics*, 12(2):1, 2011.
- [] Jared T Simpson, Kim Wong, Shaun D Jackman, Jacqueline E Schein, Steven JM Jones, and Inanç Birol. Abyss: a parallel assembler for short read sequence data. *Genome research*, 19(6):1117–1123, 2009.
- [] Yu Peng, Henry C.M. Leung, S.M. Yiu, and Francis Y.L. Chin. Idbba-ud: a de novo assembler for single-cell and metagenomic sequencing data with highly uneven depth. *Bioinformatics*, 28:1420–1428, 2012.
- [] Anton Bankevich, Sergey Nurk, Dmitry Antipov, Alexey A. Gurevich, Mikhail Dvorkin, Alexander S. Kulikov, Valery M. Lesin, Sergey I. Nikolenko, Son Pham, Andrey D. Pribelski, Alexey V. Pyshkin, Alexander V. Sirotkin, Nikolay Vyahhi, Glenn Tesler, Max A. Alekseyev, and Pavel A. Pevzner. Spades: A new genome assembly algorithm and its applications to single-cell sequencing. *Journal of Computational Biology*, 19(5):455–477, 2012.
- [] Dinghua Li, Ruibang Luo, Chi-Man Liu, Chi-Ming Leung, Hing-Fung Ting, Kunihiro Sadakane, Hiroshi Yamashita, and Tak-Wah Lam. Megahit v1.0: A fast and scalable metagenome assembler driven by advanced methodologies and community practices. *Methods*, 102:3–11, 2016.
- [] Shakya Migun, Christopher Quince, James Campbell, Zamin Yang, Christopher Schadt, and Mircea Podar. Comparative metagenomic and rRNA microbial diversity characterization using archaeal and bacterial synthetic communities. *Environmental Microbiology*, 15(6):1882–1899, 2013.
- [] H Chitsaz, JL Yee-Greenbaum, G Tesler, MJ Lombardo, CL Dupont, JH Badger, M Novotny, DB Rusch, LJ Fraser, NA Gormley, O Schulz-Trieglaff, GP Smith, DJ Evers, PA Pevzner, and RS Lasken. Efficient de novo assembly of single-cell bacterial genomes from short-read data sets. *Nat Biotechnol*, 29(10):915–21, 2011.
- [] Anthony M. Bolger, Marc Lohse, and Bjoern Usadel. Trimmomatic: A flexible trimmer for illumina sequence data. *Bioinformatics*, 30(15):2114–2120, 2014.
- [] Matthew D MacManes. On the optimal trimming of high-throughput mRNA sequence data. *Frontiers in genetics*, 5:13, 2014.
- [] Heng Li and Richard Durbin. Fast and accurate short read alignment with burrows-wheeler transform. *Bioinformatics*, 25(14):1754–1760, 2009.
- [] Heng Li, Bob Handsaker, Alec Wysoker, Tim Fennell, Jue Ruan, Nils Homer, Gabor Marth, Goncalo Abecasis, Richard Durbin, et al. The sequence alignment/map format and samtools. *Bioinformatics*, 25(16):2078–2079, 2009.
- [] Stefan Kurtz, Adam Phillippy, Arthur L Delcher, Michael Smoot, Martin Shumway, Corina Antonescu, and Steven L Salzberg. Versatile and open software for comparing large genomes. *Genome biology*, 5(2):1, 2004.




Figure 2. NGA50 by genome and assembler.