

***F1000Research* Evaluating Metagenome Assembly on a Complex Community**

Sherine Awad¹, Luiz Irber², and C. Titus Brown³

^{1,2,3}Department of Population Health and Reproduction, University of California, Davis, California

Abstract

Metagenome assembly is a challenging problem due to the biodiversity of the microorganisms. Most assemblers are designed for whole genome assembly and not capable of dealing with metagenomic samples. However, in order to decide which assembler works best for metagenome, we need to evaluate metagenome assembly generated by each assembler.

In this paper, we used three assemblers ; IDBA-UD, SPAdes, and Megahit to assemble metagenome mock community data and evaluate the assembly process in terms of resources utilization, assembly quality, genome fraction covered, duplication ratio, misassemblies and partial alignments.

The results show only small differences in content recovery between assemblers. However, Megahit is much faster and produces shorter contig lengths than IDBA-UD and SPAdes.

Introduction

Metagenomics refers to sequencing of DNA from a mixture of organisms, often from an environmental or uncultured sample. Unlike whole genome sequencing, metagenomics targets a mixture of genomes, which introduces metagenome-specific challenges in analysis. Most approaches to analyzing metagenomic data rely on mapping or comparing sequencing reads to reference sequence collections. However, reference databases contain only a small subset of microbial diversity (cite: geba), and the much of the remaining diversity is evolutionarily distant and search techniques may not access it.

As sequencing capacity increases and sequence data is generated from many more environmental samples, metagenomics is increasingly using de novo assembly techniques to generate new reference genomes and metagenomes. There are a number of metagenome assemblers that are widely used. However, evaluating the results of these assemblers is challenging due to the general lack of good quality reference metagenomes. Below, we evaluate three commonly assemblers - SPAdes, IDBA, and MEGAHIT - on a mock community containing 64 species of microbes with known genomes.

Moya et al. in [?] evaluated metagenome assembly using simulated two 454 viral metagenome and six assemblers. The assemblies were evaluated based on several metrics including N50, percentages of reads assembled, accuracy when compared to the reference genome. In addition to, chimeras per contigs and the effect of assembly on taxonomic and functional annotations.

Mavromatis et al. in [?] provided a benchmark study to evaluate the fidelity of metagenome process methods. The study used simulated metagenomic data sets constructed at different complexity levels. The datasets were assembled using Phrap v3.57, Arachne v.2 [?] and JAZZ. [?] This study evaluates assembly, gene prediction, and binning methods. However, the study did not evaluate the assembly quality against a reference genome.

Rangwala et al. in [?] presented an evaluation study of metagenome assembly. The study used a de Bruijn graph based assembler ABYSS [?] to assemble simulated metagenome reads of 36 bp. The data set is classified at different complexity levels. The study compares the quality of the assembly of the data sets in terms of quality measures of contigs length, assembly accuracy. The study also took into consideration the effect of kmer size and the degree of chimericity. However, the study evaluated the assembly based on one assembler, and did not evaluate assembly against several assemblers. Also, both previous studies used simulated data, which may lack confounders of assembly such as sequencing artifacts and GC bias.

Shakya et al. (2013) constructed a complex synthetic community of organisms by mixing DNA isolated from individual cultures of 64 bacteria and archaea. In addition to performing 16s amplicon analysis and doing 454 sequencing, the authors shotgun sequenced the mixture with Illumina (@cite). While the authors concluded that this metagenomic sequencing generally outperformed amplicon sequencing, they conducted a map-

ping based analysis rather than an assembly based analysis.

In this paper, we evaluate metagenome assembly on the Illumina data set from Shakya et al. (2013) using three assemblers; IDBA-UD [?], SPAdes [?], and MEGAHIT [?]. These three assemblers were chosen because they are actively used and highly cited.

SPAdes [?] is an assembler for both single-cell and standard (multicell) assembly. (More description here @CTB.)

IDBA-UD [?] is a de Bruijn graph approach for assembling reads from single cell sequencing or metagenomic sequencing technologies with uneven sequencing depths. IDBA-UD uses multiple depth-relative thresholds to remove erroneous k-mers in both low-depth and high-depth regions. It also uses paired-end information to solve the branch problem of low-depth short repeat regions. It also applies an error correction step to correct reads of high-depth regions that can be aligned to high confident contigs.

MEGAHIT [?] is a newer approach that constructs a succinct de Bruijn graph using multiple k-mer sizes, and uses a novel “mercy k-mer” approach that preserves low-abundance regions of reads. It also can use GPUs to accelerate the graph construction.

Below, we evaluate the performance of these three assemblers using the synthetic (“mock”) community data from the Shakya et al. study. The performance of each assembler is compared in terms of resource utilization, covered genome fraction, duplication ratio, gene recovery, contig misassembly, and contig length.

In contrast to some other evaluations (assemblathon 2), we work with a synthetic community with a known answer, do not tune the default parameters, and investigate memory and runtime as a key part of the process. This mimics default user experience. (Expand.)

Table 1. Running Time and Memory Utilization

(1) IDBA-UD	
Running Time	17:12:43
Memory Utilization (GB)	149.12
(2) SPAdes	
Running Time	42:14:06
Memory Utilization (GB)	391.45
(3) MEGAHIT	
Running Time	56:04.43
Memory Utilization (GB)	34.40

Datasets

We used a diverse mock community data set constructed by pooling DNA from 64 species of bacteria and archaea and sequencing them with Illumina HiSeq. The raw data set consisted of 109,629,496 reads from Illumina HiSeq 101 bp paired-end sequencing (2x101) with an untrimmed total length of 11.07 Gbp and an estimated fragment size of 380 bp [?].

Table 2. Reference Genome Coverage and Duplication Ratio

(1) Best hit Approach		
(1) IDBA-UD		
99.0	Genome Coverage	56.89 %
	Duplication Ratio	0.38 %
95.0	Genome Coverage	58.00%
	Duplication Ratio	0.59 %
(2) SPAdes		
99.0	Genome Coverage	63.79 %
	Duplication Ratio	0.15 %
95.0	Genome Coverage	64.68 %
	Duplication Ratio	0.26%
(3) MEGAHIT		
99.0	Genome Coverage	68.47 %
	Duplication Ratio	0.37 %
95.0	Genome Coverage	68.96 %
	Duplication Ratio	0.45%
(2) Ambiguous Approach		
(1) IDBA-UD		
99.0	Genome Coverage	89.79 %
	Duplication Ratio	0.94%
95.0	Genome Coverage	95.46 %
	Duplication Ratio	1.90 %
(2) SPAdes		
99.0	Genome Coverage	89.42 %
	Duplication Ratio	1.00 %
95.0	Genome Coverage	95.12 %
	Duplication Ratio	1.98 %
(3) MEGAHIT		
99.0	Genome Coverage	91.16 %
	Duplication Ratio	0.55 %
95.0	Genome Coverage	94.22 %
	Duplication Ratio	1.48 %
(3) No Misassemblies Approach		
(1) IDBA-UD		
99.0	Genome Coverage	34.60 %
	Duplication Ratio	0.37%
95.0	Genome Coverage	35.22 %
	Duplication Ratio	0.56 %
(2) SPAdes		
99.0	Genome Coverage	35.92%
	Duplication Ratio	0.16 %
95.0	Genome Coverage	36.42 %
	Duplication Ratio	0.21 %
(3) MEGAHIT		
99.0	Genome Coverage	45.81 %
	Duplication Ratio	0.39%
95.0	Genome Coverage	46.03 %
	Duplication Ratio	0.48 %

Table 3. Contigs Analysis

(1) Best hit Approach	
(1) IDBA-UD	
No. of Contigs	19,988
Totally Aligned Contigs %	72.96% (97,138,779)
Partial Aligned Contigs %	10.97% (20,261,669)
Unaligned Contigs %	16.07% (61,421,243)
(2) SPAdes	
No. of Contigs	15,254
Totally Aligned Contigs%	76.52% (109,342,809)
Partial Aligned Contigs%	12.08% (22,008,234)
Unaligned Contigs%	11.40% (34,176,209)
(3) MEGAHIT	
No. of Contigs	27,657
Totally Aligned Contigs %	83.07% (128,987,917)
Partial Aligned Contigs%	4.67% (12,325,804)
Unaligned Contigs%	12.26% (50,093,476)
(2) Ambiguous Approach	
(1) IDBA-UD	
No. of Contigs	19,988
Totally Aligned Contigs%	80.59% (161,075,933)
Partial Aligned Contigs%	8.54 % (22,638,415)
Unaligned Contigs %	10.87% (13,378,572)
(2) SPAdes	
No. of Contigs	15,254
Totally Aligned Contigs	81.03 % (154,920,366)
Partial Aligned Contigs%	9.34% (28,028,529)
Unaligned Contigs%	9.62% (12,931,934)
(3) MEGAHIT	
No. of Contigs	27,657%
Totally Aligned Contigs%	87.59 % (169,789,173)
Partial Aligned Contigs%	4.36% (15,658,616)
Unaligned Contigs%	8.04% (12,777,886)
(3) No Misassemblies Approach	
(1) IDBA-UD	
No. of Contigs	19,988
Totally Aligned Contigs%	57.41% (61,874,288)
Partial Aligned Contigs%	6.81%(9,530,232)
Unaligned Contigs %	35.78% (122,899,406)
(2) SPAdes	
No. of Contigs	15,254
Totally Aligned Contigs	62.93% (64,656,150)
Partial Aligned Contigs%	6.84% (9,323,121)
Unaligned Contigs%	30.23% (114,292,858)
(3) MEGAHIT	
No. of Contigs	27,657
Totally Aligned Contigs%	67.62% (87,488,973)
Partial Aligned Contigs%	2.73% (7,069,258)
Unaligned Contigs%	29.65% (102,672,613)

The original reads are available through the NCBI Sequence Read Archive at Accession SRX200676. We received the 64 reference genomes from the original

authors. They consist of 205.6 Mbp of assembled genomes in 64 contigs, and are available for download at <https://dx.doi.org/10.6084/m9.figshare.1506873.v2>

Table 4. Read Mapping Analysis

Assembler	Read that map	Read that map uniquely
IDBA-UD	7,067,517	2,332,735
SPAdes	6,786,747	1,966,038
Megahit	6,392,926	1,585,363

Table 5. Genomes with the most common uncovered bases between the three assemblies.

Genome	Uncovered bases (% of total)
Shewanella baltica OS223	2.3 Mbp (18.25%)
Fusobacterium nucleatum	2.2 Mbp (16.95%)
Desulfovibrio vulgaris DP4	1.6 Mbp (12.85%)
Enterococcus faecalis V583	1.0 Mbp (7.83%)
Thermus thermophilus HB27	0.7 Mbp (5.73%)

Table 6. Analysis of top Common Uncovered Genomes

Genome	Fraction Assembled	Uncovered by Reads
Shewanella baltica OS223	(iqc 23.35%), (sqc 21.80%), (mqc 47.41%)	0.05%
Fusobacterium nucleatum	(iqc 0%), (sqc 0%), (mqc 0.18%)	52.70%
Desulfovibrio vulgaris DP4	(iqc 49.38%) (sqc 48.83%) (mqc 50.01%)	8.02%
Enterococcus faecalis V583	(iqc 67.07%) (sqc 66.53%) (mqc 68.50%)	21.85%
Thermus thermophilus HB27	(iqc 59.85%) (sqc 58.74%) (mqc 58.44%)	8.93%

Table 7. Analysis of reads and kmer abundance

Assembler	Unmapped Reads	High abundance k-mers not in assemblies
IDBA-UD	3,328,674	
SPAdes	3,879,573	
Megahit	5,848,494	

Methods

The analysis code and run scripts for this paper are available at: <https://github.com/dib-lab/2015-metagenome-assembly/>. The scripts and overall pipeline were examined by the first and senior authors for correctness. In addition, the bespoke reference-based analysis scripts were tested by running them on a single-colony *E. coli* MG1655 data set which had a high quality reference genome (cite Chitsaz, 2011).

Table 8. Misassembled Contigs using Identity 99%

Assembler	No. of Contigs	No. of Bases
IDBA-UD	4,980 (24.91%)	112,309,828
SPAdes	3,143 (20.60%)	108,969,624
Megahit	5,977 (21.61%)	90,889,558




Figure 1. Cumulative coverage profile for the reference metagenome, based on read mapping.

Quality Filtering

We removed adapters with Trimmomatic v0.30 in paired-end mode with the Truseq adapters [?], using light quality score trimming as recommended in @cite MacManes 2014.

Reference Coverage Profile

To evaluate how much of the reference metagenome was contained in the read data, we used `bwa aln` to map reads to the reference genome. We then calculated how many reference bases were covered by how many mapped reads (custom script `coverage-profile.py`).

Assemblers

We assembled the quality-filtered reads using three different assemblers: IDBA-UD [?], SPAdes [?], and MEGAHIT [?]. For IDBA-UD v1.1.1 [?], we used `-pre_correction` to perform pre-correction before assembly and `-r` for the `pe` files.

For SPAdes v3.9.0 [?], we used `-meta -pe1-12 -pe1-s` where `-meta` is recommended when working with metagenomic data sets, `-pe1-12` specifies the interleaved reads for the first paired-end library, and `-pe1-s` provides the orphan reads remaining from quality trimming.

For MEGAHIT [?], we used `-l 101 -m 3e9 -cpu-only` where `-l` is for maximum read length, `-m` is for max memory in bytes to be used in constructing the graph, and `-cpu-only` to use only the CPU and not the GPU. We also used `-presets meta-large` for large and complex metagenomes, and `-12` and `-r` are parameters that specify the interleaved-paired-end and single-end files respectively.

All three assemblies were executed on the same high-memory buy-in node on the Michigan State University

High Performance Compute Cluster, and we recorded RAM and CPU time of each assembly job using the `qstat` utility at the end of each run.

Unless otherwise mentioned, we eliminated all contigs less than 500 bp from each assembly prior to further analysis.

Mapping

We aligned all quality-filtered reads to the reference metagenome with `bwa aln` (v0.7.7.r441) [?]. We aligned paired-end and orphaned reads separately using `bwa aln samse`. We then used `samtools` (v0.1.19) [?] to convert SAM files to BAM files for both paired-end and orphaned reads. To count the unaligned reads, we included only those records with the “4” flag in the SAM files [?].

To extract the reads that contribute to unaligned contigs, we mapped the quality filtered reads to the unaligned contigs using `bwa aln` (v0.7.7.r441) [?]. Then we used `samtools` to retrieve the reads that mapped to the unaligned contigs.

k-mer Presence

In order to examine k-mer presence for a k-mer size of 20, we built a k-mer counting table from the given quality filtered reads using `load-into-counting.py` from `khmer` (@cite khmer). Then we calculate abundance distribution of the k-mers in the quality filtered reads using the pre-made k-mer counting table using `abundance-dist.py`. We followed the same approach to examine k-mer presence in assemblies.

Assembly analysis using Nucmer

We used the `Nucmer` tool from `MUMmer3.23` [?] to align assemblies to the reference genome with options `-coords -p`. Then we parsed the generated “coords” file using a custom script `analyze_assembly.py` to calculate several analysis metrics across all the assemblies at two alignment identities, 95% and 99%.

Reference-based analysis of the assemblies

We processed the alignments in three different ways: ambiguous, best-hit, and no-misassemblies.

In the ambiguous approach, we took into account all the alignments of a contig to the reference.

In the best-hit approach, among all alignments of a contig, we took into consideration only the alignment with the best score.

In the no-misassemblies approach, we only counted contigs that have precisely one alignment to the reference.

In all approaches, we flag a base in the reference genome as “covered” if it is contained in a kept alignment. We define the duplication ratio as the percentages of bases in the reference covered by two or more kept alignments. We define misassemblies as those contigs that are divided into different parts when mapped to the reference. The number of misassembled contigs is equal to the number

of aligned contigs (both totally and partially) in the ambiguous approach, minus the number of aligned contigs in the no-misassemblies approach.

All approaches have a non-zero duplication ratio within the reference because we do not explicitly discard contigs that map to the same location in the reference.

Gene annotations using Prokka

We used `prokka` [?] to annotate the reference genome using `-outdir mprokka -prefix testasm -metagenome`. Then we parsed the `testasm.tbl` output file to get the coordinates of CDS genes. We then searched the alignments for how many genes were contained in those alignments.

Results

The raw data is high quality

The reads contains 5,536,289,548 bp (5.536 Gbp) in 54,814,748 paired reads with 101.0 average length in total 11,072,579,096 bp (11.07 Gbp) in 109,629,496 reads with 101.0 average length.

After trimming, we got 108,422,358 paired reads with 10,938,648,751 bps (10.94 Gbp) and 100.9 average length. We also got 520,403 orphan reads with 46,563,015 bp (46.56 Mbp) with 89.5 average length. In total, 108,942,761 reads 10,985,211,766 bp (10.98 Gbp) with 100.8 average length. Trimming removed 686,735 reads (0.63 %) of the data, which means the data set is of high quality. This quality trimmed (“QC”) data set was used as the basis for all further analyses.

98% or more of the reference is present in the read data set

We next evaluated the of the reference genome covered by at least one read (see Methods for details). Quality filtered reads cover 203,058,414.0 (98.76%) bases of the reference metagenome (205,603,715 bp total size). Figure 1 shows the cumulative coverage profile of the reference metagenome, and the percentage of bases with that coverage. Most of the reference metagenome was covered at least minimally; only 3.33% of the reference metagenome had mapping coverage <5, and 1.24% of the bases in the reference were not covered by any reads in the QC data set.

In order to evaluate reconstructability with De Bruijn graph assemblers, we next examined k-mer presence for a k-mer size of 20. Of the 174m 20-mers in the reference metagenome, 98.7% were present in the QC reads and 95.5% of them occurred with abundance 5 or greater in the QC reads.

7.6 million reads (XX%) from the QC data set did not map anywhere in the reference.

MEGAHIT is the fastest and lowest-memory assembler evaluated

We ran three commonly used metagenome assemblers on the QC data set: IDBA-UD, SPAdes, and Megahit. We

recorded the time and memory usage of each (Table 1). MEGAHIT outperformed both SPAdes and IDBA-UD considerably, producing an assembly in one hour – approximately 17 times faster than IDBA and 42 times faster than SPAdes. MEGAHIT used only 34.4 GB of RAM – 1/5 to 1/11th the memory used by IDBA and SPAdes, respectively.

The assemblies contain most of the raw data

We assessed read inclusion in assemblies by mapping the QC reads to the assemblies and counting the remaining unmapped reads. Depending on the assembly, between 3.3 million and 5.8 million reads (3-6% XX) did not map to the assembly. (@CTB Put in table.)

Much of the reference is covered by the assemblies

We next evaluated the extent to which the assembled contigs recovered the “known/true” metagenome sequence by aligning each assembly to the reference (Table 2). All three assemblers generate contigs that cover more than 89% of the reference metagenome at high identity (99%) with little duplication (0.55 -1.0%) (see “Ambiguous approach” in Table 2). If we relax the identity threshold to 95%, then the assembled contigs cover more than 94% of the reference metagenome.

When we use only the highest-scoring alignment at 99% identity, we find that the reference coverage drops to 57-68%, depending on the assembler (“Best hit approach”, Table 2). The reference coverage from highest-scoring alignments does not substantially increase when the identity threshold is relaxed to 95%.

At 99% identity with the ambiguous approach, approximately 6.23% of the reference is covered by no contig from any of the three assemblies; we discuss this in more detail below.

The generated contigs are broadly accurate

When counting only the best alignment per contig at a 99% identity threshold, more than 72% of contigs align to the reference completely, i.e. across the whole length of the contig (Table 3, “Best hit”, Totally Aligned Contigs %). If we allow multiple alignments per contig, then more than 80% of contigs align completely (although not contiguously) to the reference (Table 3, “Ambiguous”, Totally Aligned Contigs %). Approximately 4 - 9% of the contigs align only partially, and the remaining 8 - 10% of contigs do not align at all to the reference (discussed in detail below).

All three assemblers perform equally well in reference base recovery

86.13% of the reference metagenome is recovered in at least one contig by all three assemblers (“common covered”) with relatively little duplication, under 99% identity/ambiguous matching: 6.23% of the reference is missed by all three assemblers, 0.64% is recovered

by IDBA alone, 0.74% is recovered by SPAdes only, and 1.92% is recovered by MEGAHIT only.

Of the uncovered bases in common between the assemblers, 19.66% have zero mapping coverage, while 29.20% have mapping coverage <5.

IDBA, SPAdes, and MEGAHIT each uniquely fail to recover about 1% of the reference.

Large portions of several reference genomes are not assembled by any assembler

A number of the genomes in the reference metagenome had many missing bases in the assemblies (Table 5). In three extreme cases, *Shewanella baltica* OS223, *Fusobacterium nucleatum*, and *Desulfovibrio vulgaris* contribute 18.25%, 16.95%, and 12.85% of the total common uncovered bases respectively.

Most genes within the reference metagenome are contained within contigs

The reference genome has 188,880 CDS with 91,806 annotated as genes, based on a Prokka annotation (see Methods). Using the “ambiguous” alignment approach and 99% identity, the IDBA-UD assembly contained 82,791 (95.20%) of the reference genes, while SPAdes contained 83,475 (94.44%), and MEGAHIT contained 80,256 (94.59%) of the reference genes.

Many assembled contigs do not align to the reference metagenome

Depending on assembler, between 8.04% and 10.87% of the assembled contigs (XX-YY bases total) do not align anywhere in the reference metagenome. We extracted these contigs and analyzed their content and read mapping separately.

Between 5.9 and 6.5% of the QC reads mapped to the unalignable contigs, depending on assembler (@CTB put in table). In each case, between 67% and 75% the reads that mapped to these unalignable contigs did not map anywhere in the reference (@CTB put in table); about 5m QC reads total map only to the unaligned contigs and nowhere in the reference.

For each assembly, approximately 5m quality-filtered reads map only to the unaligned contigs (and nowhere in the reference).

We used sourmash (cite) to search the unalignable contigs against all 52,000 microbial genomes currently available from NCBI and found nine genomes with significant estimated overlap with the unalignable contigs (10% or more estimated overlap - @CTB put in Methods). We then aligned the unalignable contigs against these nine genomes, and found that between 16% and 28% of the previously unaligned contigs matched to the new genomes at 99% identity.

Discussion

Assembly recovers basic content well

The majority of the reference metagenome is recovered by all three assemblers: 89% or more of the reference metagenome is contained within each assembler's output at 99% identity, and 94% can be recovered if we relax the identity threshold to 95%. This is close to the measured maximum reconstructability based on read mapping and k-mer presence: 98.76% of the reference is covered by at least one read, and 98.7% of the genome is present in k-mers of size 20.

The contigs generated align well to the reference metagenome: with all assemblers, 80% (or more) of contigs align to the reference metagenome across the whole length of the contig, and another 4% (or more) of the contigs are entirely contained within the reference metagenome, although they do not align to only one location in the reference.

More than 6.23% of the reference metagenome is missing from all of the assemblies when 99% similarity is required, and large portions of the missing sequence are from a few genomes – in some cases, close to a fifth of the source genome is missing from the assemblies. About a third of the missing reference (29.2%) is low coverage based on read mapping, but the remainder has read coverage greater than 5. (Chase this down descriptively.)

@CTB here mention gene recovery/prokka etc is good.

Assembly produces contiguous sequence not in the reference metagenome

In addition to recovering most of the content of the reference metagenome, all three assemblers generate many contigs (approximately 10% of the total) that do not map to anything in the reference. While some could be misassemblies, many of the reads that map to these contigs do not map anywhere in the reference, while many of the contigs have coverage >XX, suggesting that they are present within the source DNA at an abundance that is similar to many of the genomes in the mock community. (@CTB identify etc.)

Whatever their true identity, these contigs are not present in the reference metagenome. Because this is a mock metagenome for which isolates were grown and individually extracted prior to being combined for sequencing, we believe that these extra genomic contigs must come from one or more contaminants. We cannot rule out contamination from kits, but because large amounts of DNA were used to create the mock community, it seems unlikely that these are due to trace contaminants created by PCR (as can happen in amplicon studies); they are most likely from contaminants grown together with the isolates.

Different assemblers differ in content recovery, but not by much

The three assemblers differ very little in recovery of basic content, as judged by reference metagenome alignments: at 99% identity, XX% is recovered in common, with 0.64 to 1.92% specific to each assembler (and a total of XX% being recovered by one or two, but not all three,

of the assemblers). None of the assemblers include all of the sequence produced by the others, suggesting that assembler-specific heuristics are responsible for the variation.

Genome assembly metrics are misleading

- content recovery is of first importance IMO
- contig n50/ng50 comparison?
- longer contigs lead to more misassemblies; suggest avoiding until long reads work well for metagenomes
- but long reads don't yet work well for metagenomes due to diversity
- fosmid type approaches maybe? moleculo etc.

Recommendation: try MEGAHIT first

We do not see a significant advantage to one assembler over another in terms of recovering the known reference, and we speculate that there is little advantage to combining assemblies from multiple assemblers, since very little extra content will be recovered.

While the other assemblers produce longer contigs, we are concerned by misassembly rates, which affect longer contigs disproportionately. Moreover, misassemblies are hard to detect without a reference. (Talk more about misassemblies.)

For these reasons, our recommendation is to use MEGAHIT with short-read data: it is significantly faster and lower memory than SPAdes and IDBA-UD, recovers essentially the same content, and – for short read data sets – does not offer significant disadvantages in terms of contig length.

(Here note that megahit does not allow long reads, unlike SPAdes and maybe IDBA. @CTB verify this.)

Conclusions

Assembly works well. There is no big difference between assemblers' performance in terms of assembly quality. In terms of cost, Megahit is much faster and utilizes less memory.

- If you have short-read sequencing only, use MEGAHIT.
- MEGAHIT can be useful for exploratory analysis of content as well.
- Metagenome assembly is good at recovering content of well covered regions. It is less good at long-range contiguity. It is better than reference based analysis at recovering content.
- Resequencing mock community genomes, as in evolutionary studies.

Author contributions

In order to give appropriate credit to each author of an article, the individual contributions of each author to the manuscript should be detailed in this section. We recommend using author initials and then stating briefly how they contributed.

Competing interests

No competing interest to our knowledge.

Grant information

This work is funded by Moore and NIH.

Acknowledgments

Michael R. Crusoe